



Universität Stuttgart

KI – Institute for Artificial Intelligence

Analytic Computing

Machine Learning

10 Support Vector Machines

Prof. Dr. Steffen Staab

Nadeen Fatallah

Daniel Frank

Akram Sadat Hosseini

Jiaxin Pan

Osama Mohamed

Arvinhd Arunbabu

Tim Schneider

Yi Wang



<https://www.ki.uni-stuttgart.de/>

- based on slides by
 - Thomas Gottron, U. Koblenz-Landau,
<https://west.uni-koblenz.de/de/studying/courses/ws1718/machine-learning-and-data-mining-1>
 - Andrew Zisserman, <http://www.robots.ox.ac.uk/~az/lectures/ml/lect2.pdf>



1 Perceptron Algorithm

Binary classification

- Given training data $\{(x_i, y_i)\}_{i=1}^N$ with $x_i \in \mathbb{R}^d$ and $y_i \in \{-1, 1\}$,
- Learn a classifier

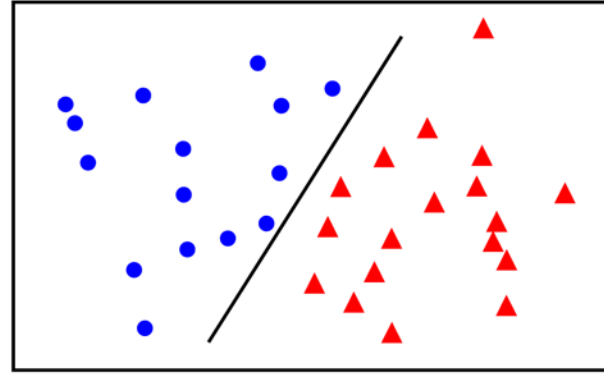
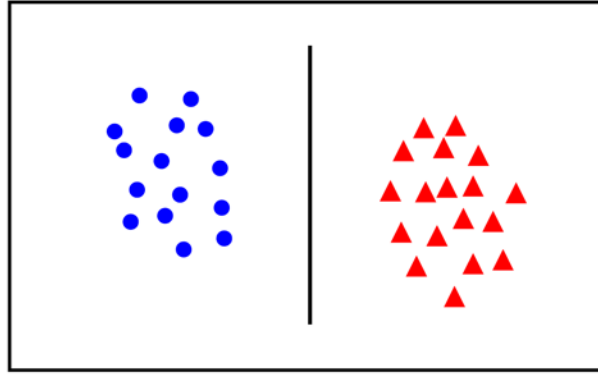
$$\hat{f}(x_i) = \begin{cases} > 0, & \text{if } y_i = +1 \\ < 0, & \text{if } y_i = -1 \end{cases}$$

- Correct classification:

$$\hat{f}(x_i)y_i > 0$$

Linear separability

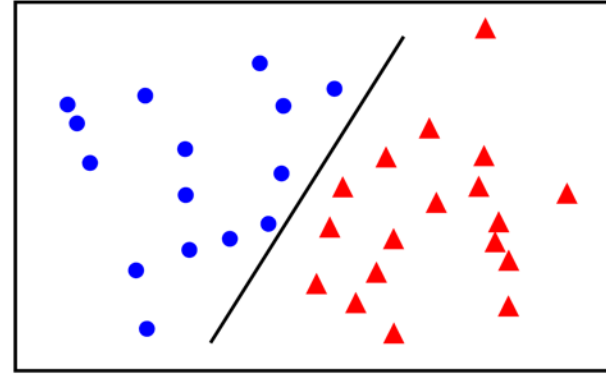
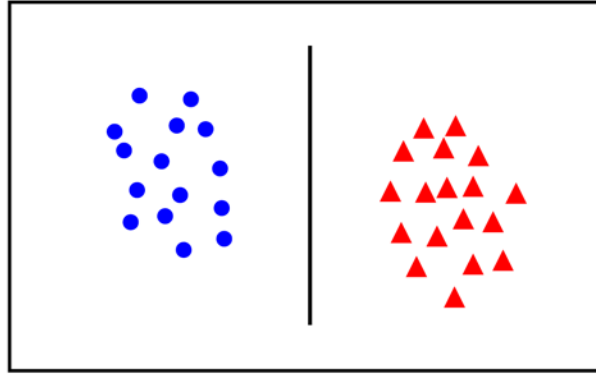
linearly
separable



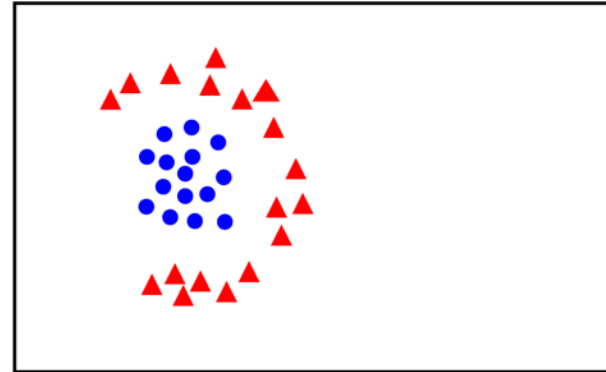
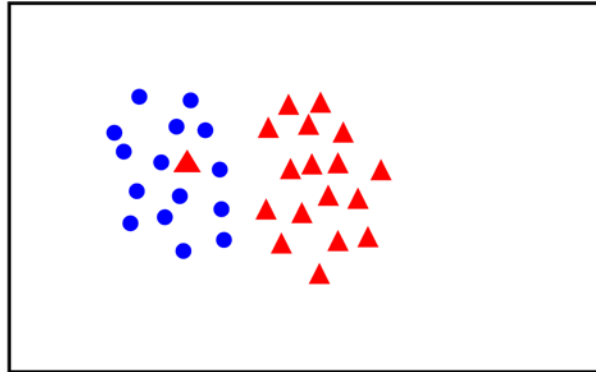
(Zisserman 2015)

Linear separability

linearly
separable



not
linearly
separable



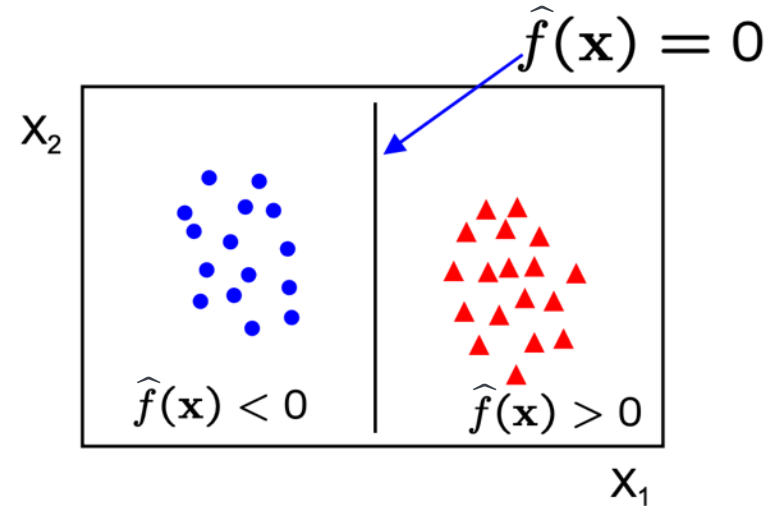
(Zisserman 2015)

Linear classifiers

- A linear classifier has the form

$$\hat{f}(x) = w^T x + b$$

- In 2D the discriminant is a line
- w is the **normal** to the line,
and b the **bias**
- w is known as the **weight vector**

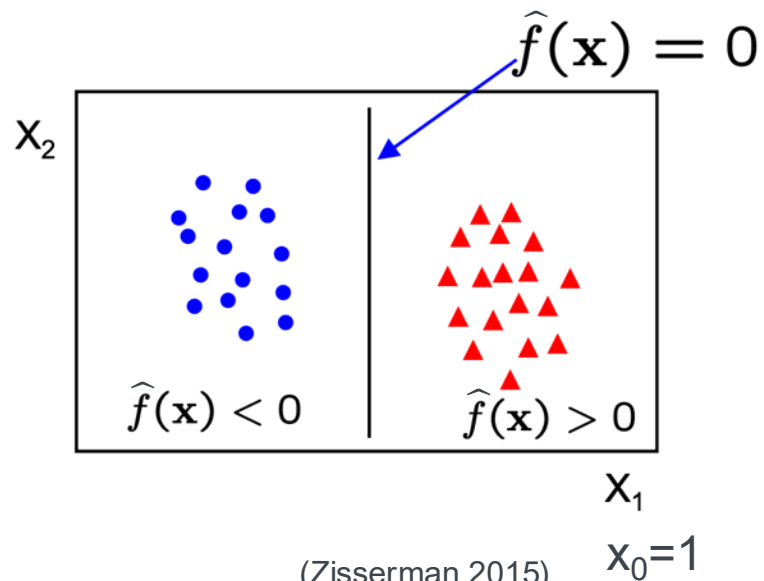


(Zisserman 2015)

Linear classifiers

- Let's assume $\mathbf{x}_i = (\mathbf{1}, x_{i,1}, \dots, x_{i,d})$
(as in linear regression)
- The we write

$$\hat{f}(x) = w^T x$$

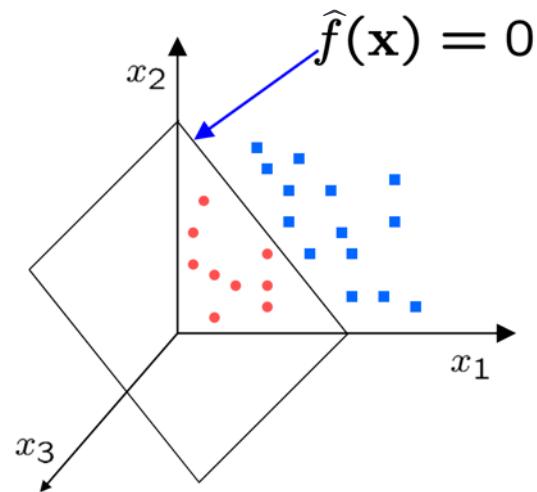


Linear classifiers

- A linear classifier has the form

$$\hat{f}(x) = w^T x$$

- In 3D the discriminant is a plane
- In n-dim the discriminant is a hyperplane
- Only w (including b) are needed to classify new data



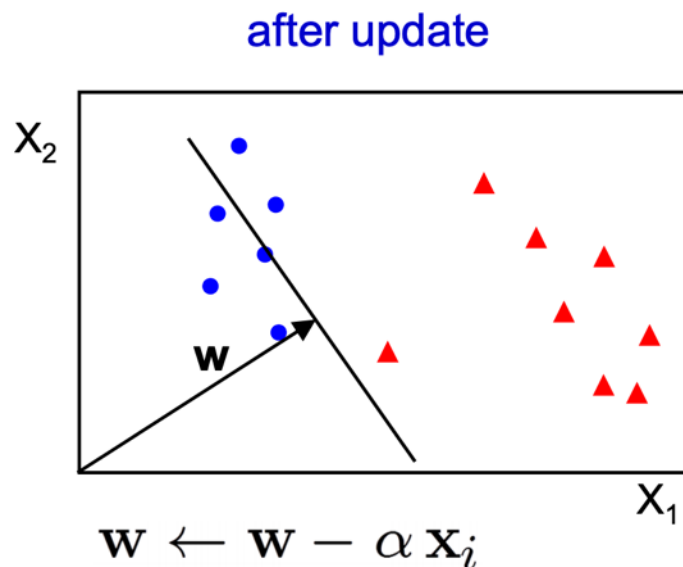
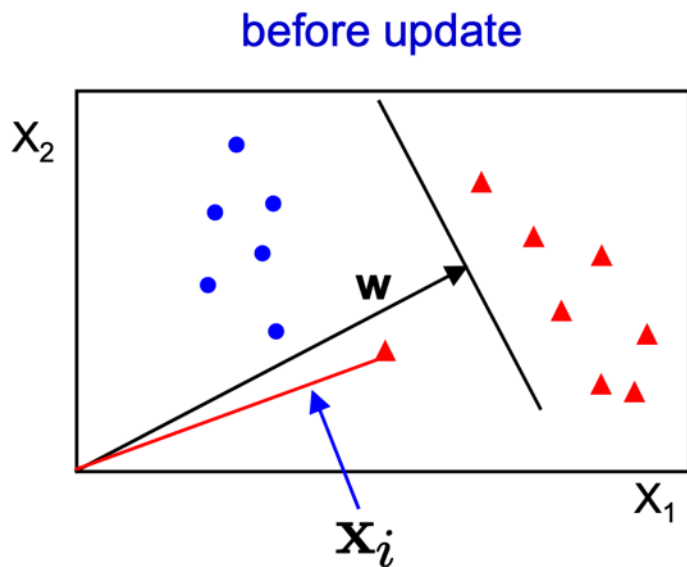
(Zisserman 2015)

The perceptron classifier

- Given training data $\{(x_i, y_i)\}_{i=1}^N$ with $x_i \in \mathbb{R}^d$ and $y_i \in \{-1, 1\}$,
- how to find a weight vector w , the *separating hyperplane*, such that the two categories are separated for the dataset?
- Perceptron algorithm
 1. Initialize $w = \bar{0}$
 2. While there is i such that $\hat{f}(x_i)y_i < 0$ do
 - $w := w - \alpha x_i \operatorname{sign}(\hat{f}(x_i)) = w + \alpha x_i y_i$ (not for $\operatorname{sign}(\dots)=0$)

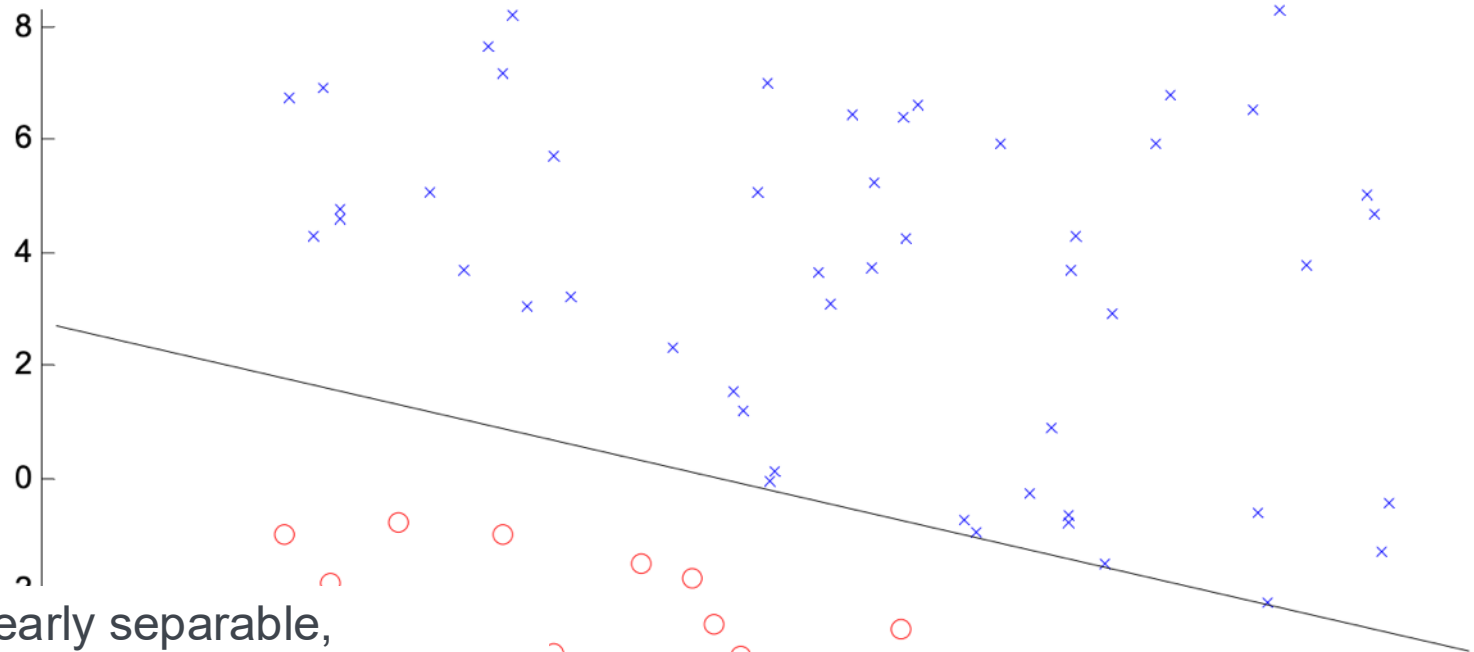
Example for perceptron algorithm in 2D

1. Initialize $w = \bar{0}$
2. While there is i such that $\hat{f}(x_i)y_i < 0$ do
 - $w := w + \alpha x_i \text{sign}(\hat{f}(x_i))$

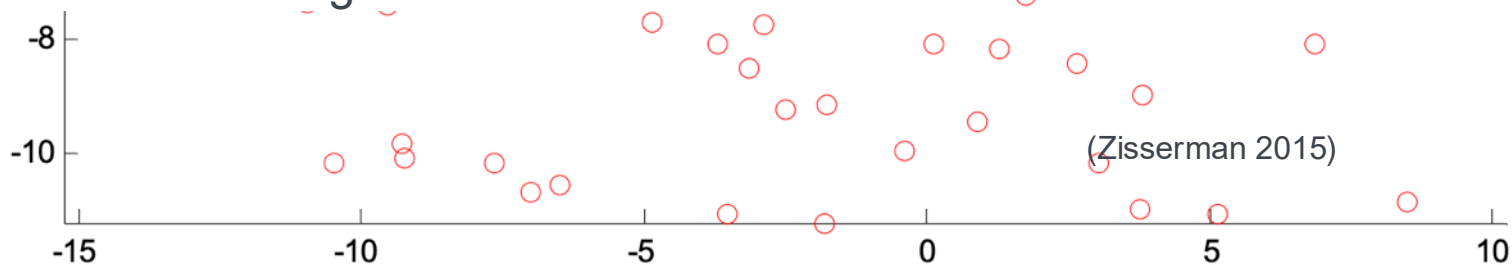


At convergence: $w = \sum_{i=1}^N \alpha_i x_i$

Example



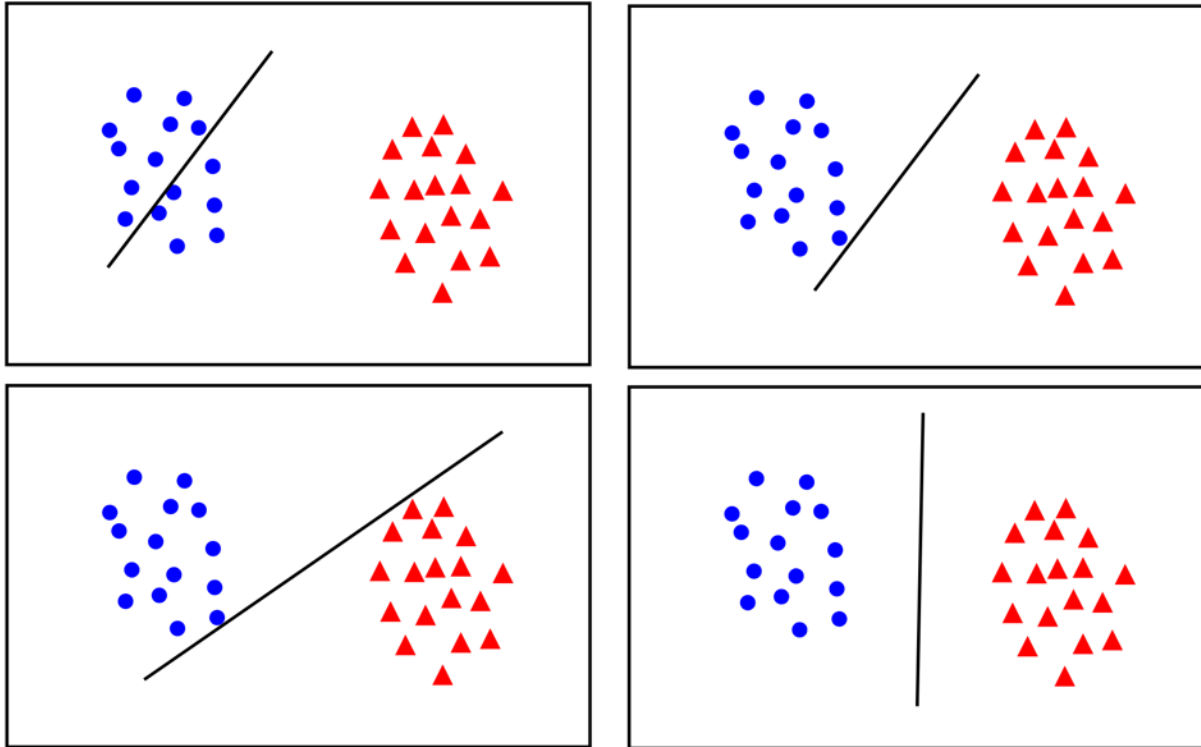
- if the data is linearly separable, then the algorithm will converge
- convergence can be slow ...
- separating line close to training data



2 Support Vectors

What is the best w ?

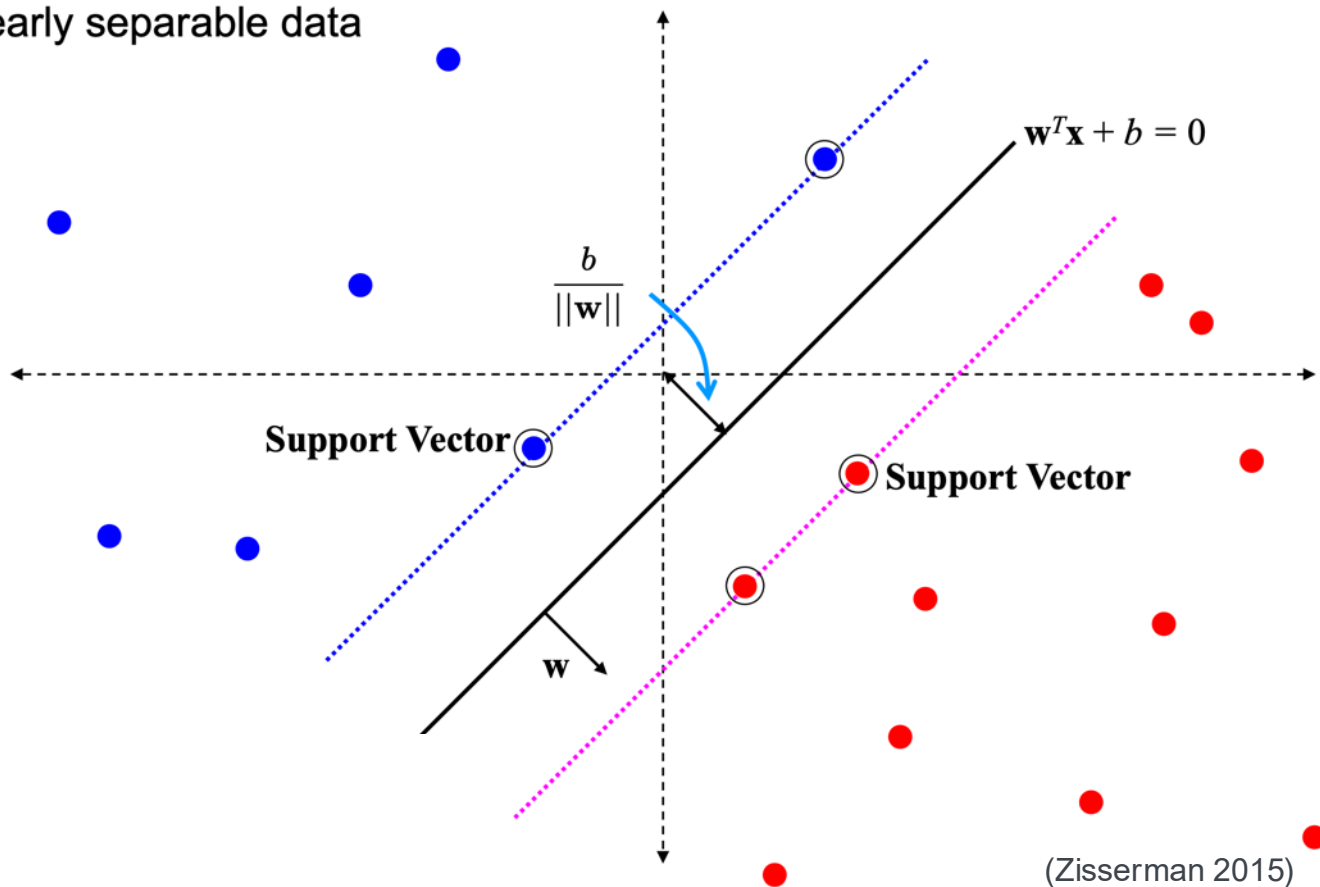
- Idea: maximum margin solution is most stable under perturbations of the inputs



(Zisserman 2015)

Support Vector Machine

linearly separable data

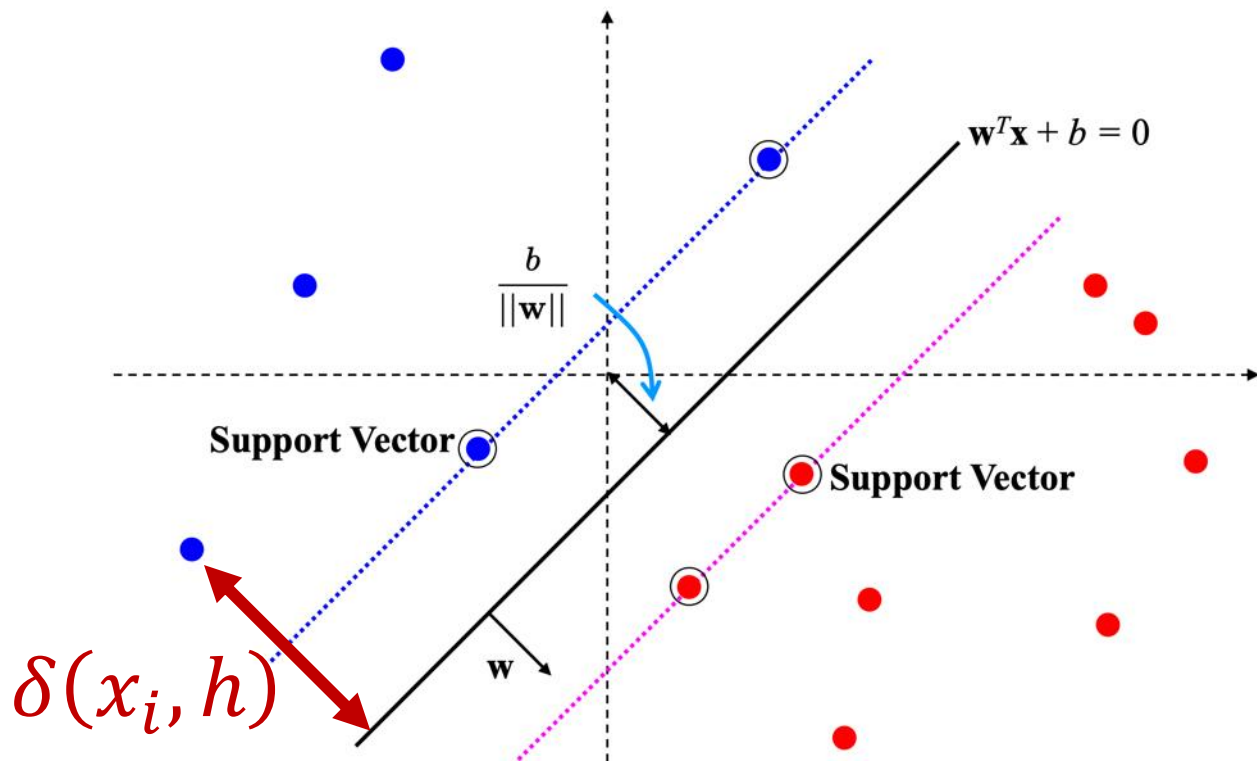


(Zisserman 2015)

SVM Optimization Problem (1)

- Distance $\delta(x_i, h)$ of data point x_i from hyperplane h

$$\delta(x_i, h) = \frac{y_i \cdot (w^T x_i + b)}{|w|}$$



SVM Optimization Problem (1)

- In general, length of vector w does not matter
 - fix w such that support vectors x_j make $y_j \cdot (w^T x_j + b) = 1$

- Then positive and negative support vectors have

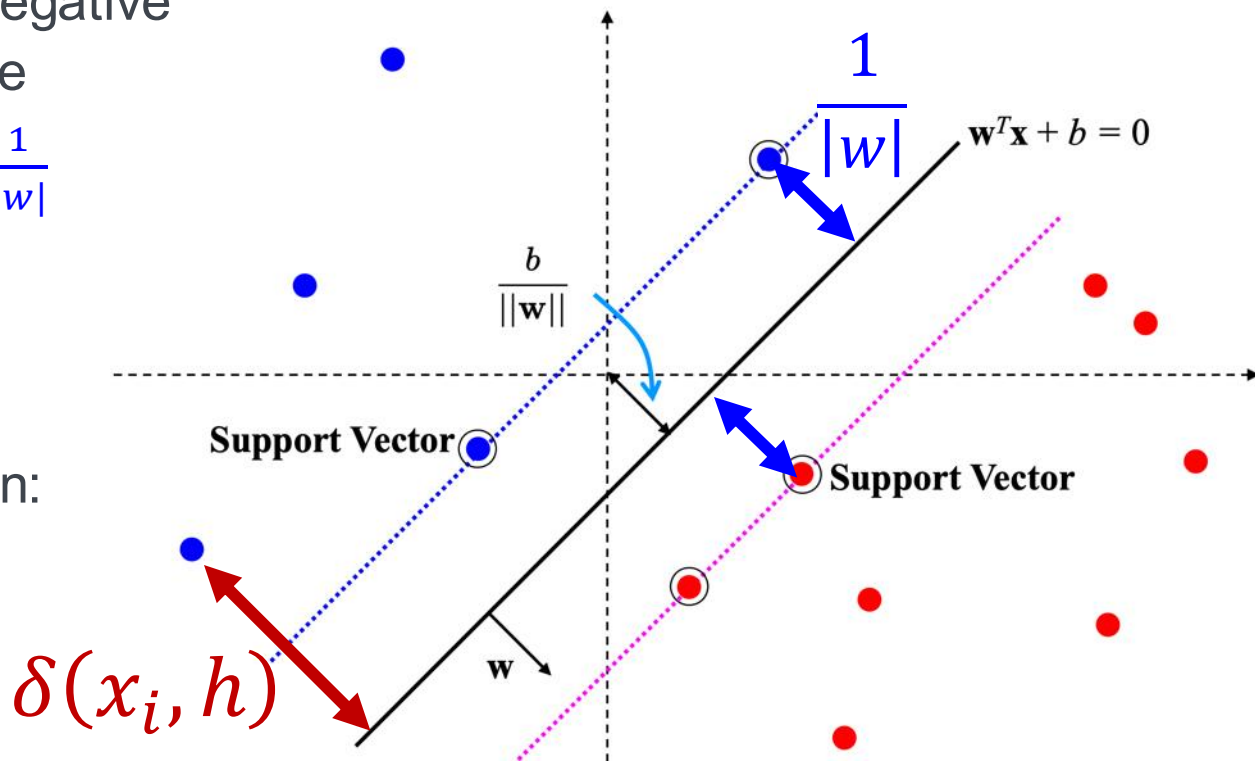
$$\text{distance } \delta(x_j, h) = \frac{1}{|w|}$$

from hyperplane –
which we want
to maximize

- Standard formulation:

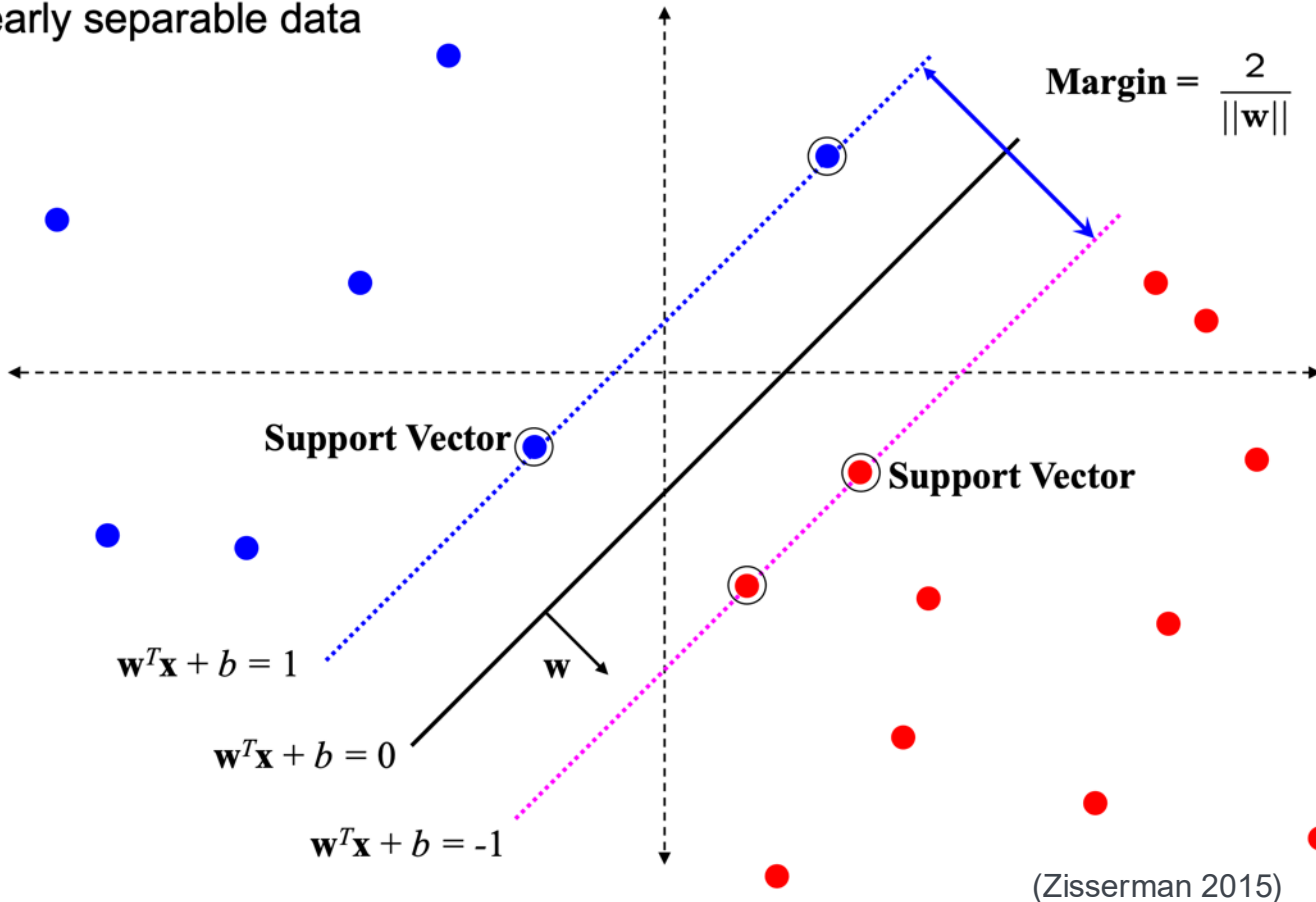
$$\text{Minimize } \frac{|w|}{2}$$

(inverse margin)



Support Vector Machine

linearly separable data



SVM Optimization Problem (2)

$$\max_w \frac{2}{||w||}$$

subject to

$$y_i(w^T x_i + b) \geq 1, \text{ for } i = 1 \dots N$$

Or equivalently

$$\min_w ||w||^2$$

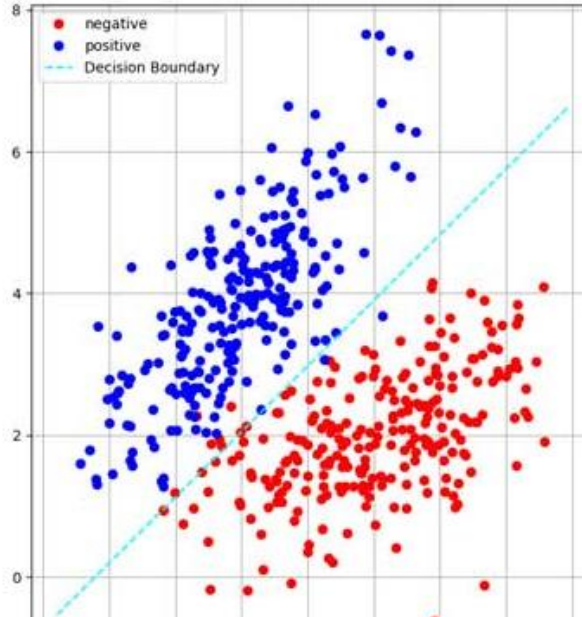
subject to the same constraints

This is a quadratic optimization problem

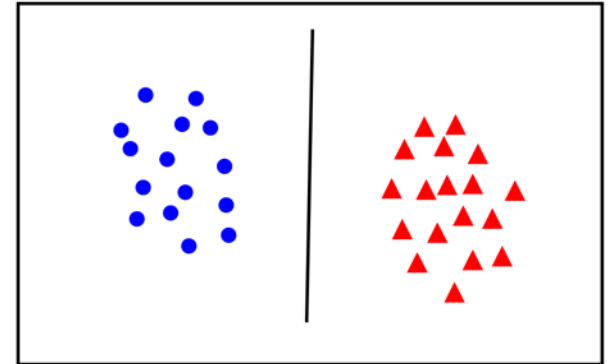
subject to linear constraints and there is a unique minimum

Compare the two optimization criteria for classification of linearly separable data by linear regression with classification by SVM

Linear classification
using regression:
decision line is average
between regression lines;
all data points are
considered



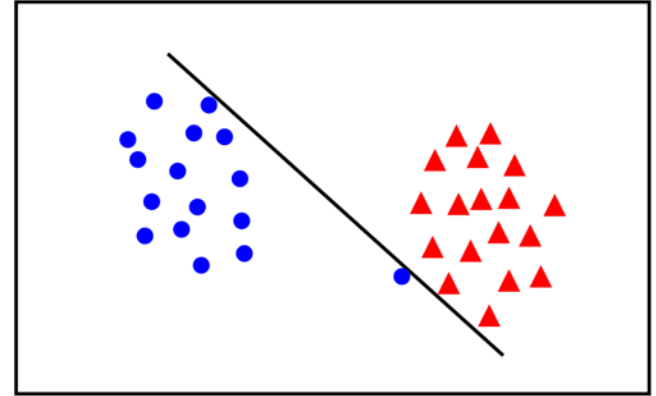
Linear classification
using SVM:
decision line maximizes
margins between support
vectors; far away data
points are irrelevant



3 Soft Margin and Hinge Loss

Re-visiting linear separability

- Points can be linearly separated, but with very narrow margin

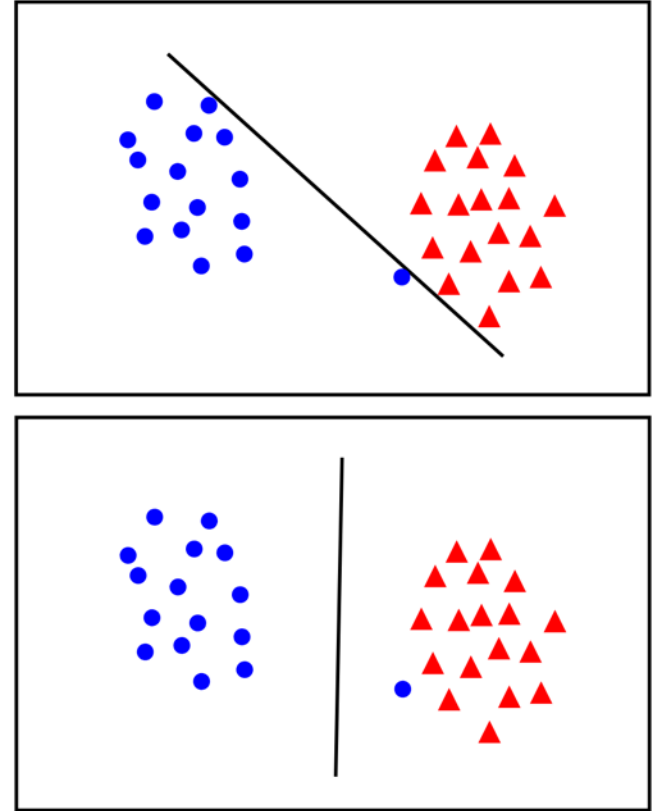


(Zisserman 2015)

Re-visiting linear separability

- Points can be linearly separated, but with very narrow margin
- Possibly the large margin solution is better, even though one constraint is violated

Trade-off between the margin and the number of mistakes on training data

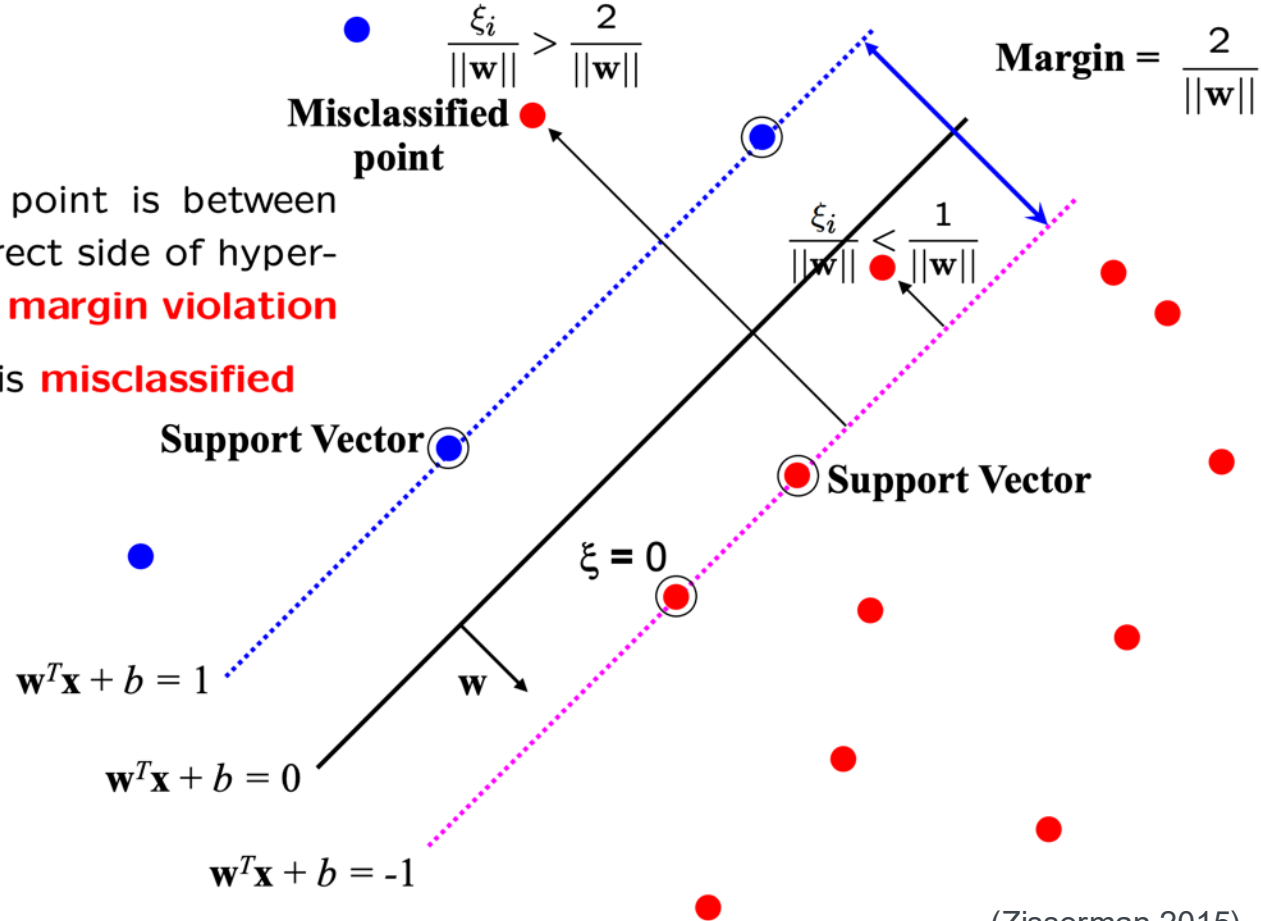


(Zisserman 2015)

Introduce „slack“ variables

$$\xi_i \geq 0$$

- for $0 < \xi \leq 1$ point is between margin and correct side of hyper-plane. This is a **margin violation**
- for $\xi > 1$ point is **misclassified**



Soft margin solution

Revised optimization problem

$$\min_{w \in \mathbb{R}^d, \xi_i \in \mathbb{R}^+} \|w\|^2 + C \sum_{i=1}^N \xi_i$$

subject to

$$y_i(w^T x_i + b) \geq 1 - \xi_i, \text{ for } i = 1 \dots N$$

- Every constraint can be satisfied if ξ_i is sufficiently large
- C is a **regularization** parameter:
 - small C allows constraints to be easily ignored \Rightarrow large margin
 - large C makes constraints hard to ignore \Rightarrow narrow margin
 - $C = \infty$ enforces all constraints \Rightarrow hard margin
- Still a quadratic optimization problem with unique minimum
- One hyperparameter C

Loss function

- Given constraints:

$$\begin{aligned}y_i(w^T x_i + b) &\geq 1 - \xi_i \\ \xi_i &\geq 0\end{aligned}$$

- We can rewrite ξ_i as:

$$\xi_i = \max(0, 1 - y_i \hat{f}(x_i))$$

- Hence, we can optimize the unconstrained optimization problem over w :

$$\min_{w \in \mathbb{R}^d} \underbrace{\frac{1}{C} \|w\|^2}_{\text{regularization}} + \sum_{i=1}^N \underbrace{\max(0, 1 - y_i \hat{f}(x_i))}_{\text{loss function}}$$

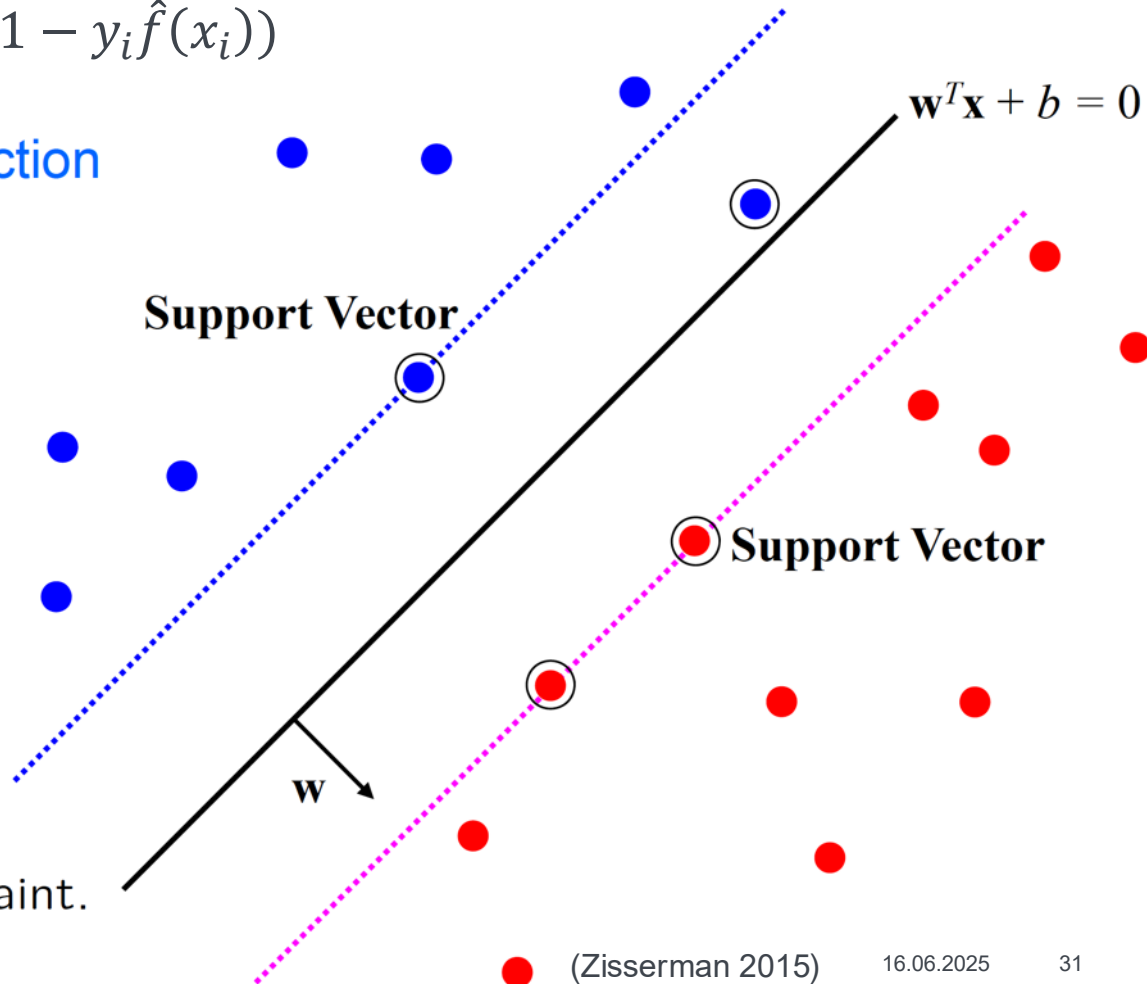
Loss function

$$\min_{w \in \mathbb{R}^d, \xi_i \in \mathbb{R}^+} \frac{1}{C} \|w\|^2 + \sum_{i=1}^N \max(0, 1 - y_i \hat{f}(x_i))$$

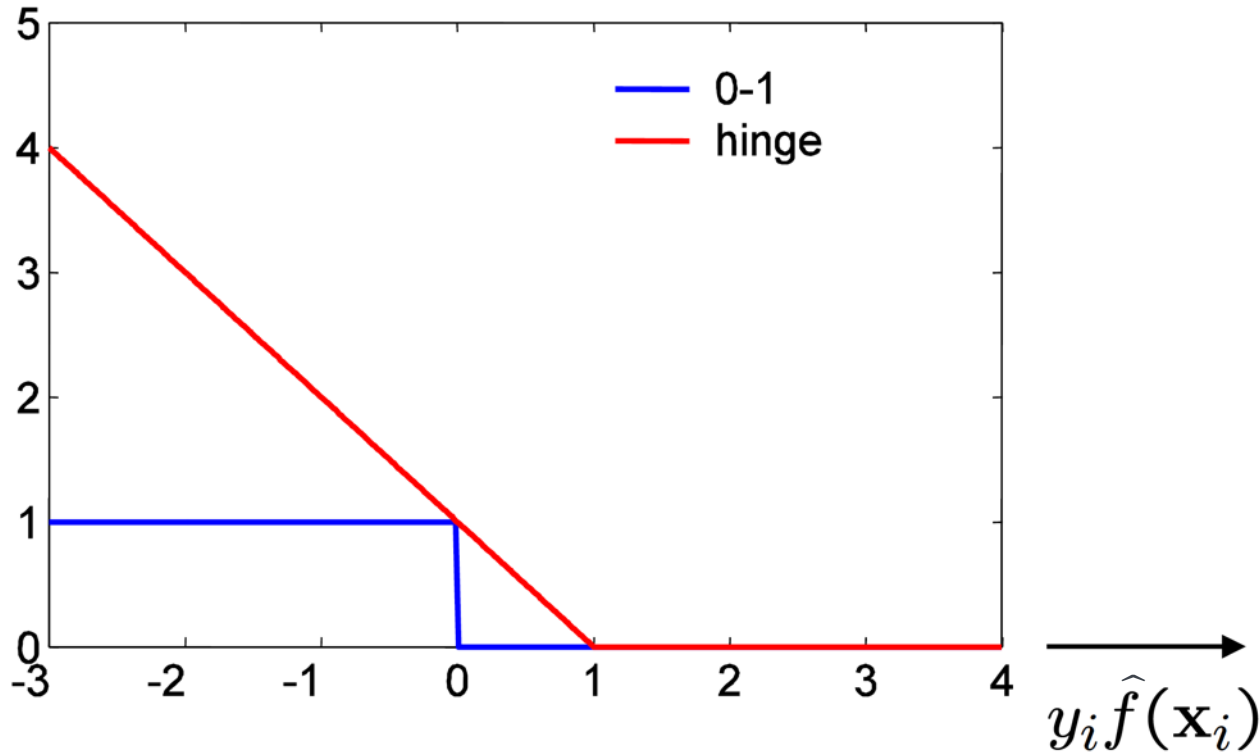
loss function

Points are in three categories:

1. $y_i f(x_i) > 1$
Point is outside margin.
No contribution to loss
2. $y_i f(x_i) = 1$
Point is on margin.
No contribution to loss.
As in hard margin case.
3. $y_i f(x_i) < 1$
Point violates margin constraint.
Contributes to loss



Hinge loss



- SVM uses “hinge” loss $\max(0, 1 - y_i \hat{f}(\mathbf{x}_i))$
- an approximation to the 0-1 loss

(Zisserman 2015)

4 Gradient descent over convex function

Gradient descent/ascent

Climb down a hill

Climb up a hill

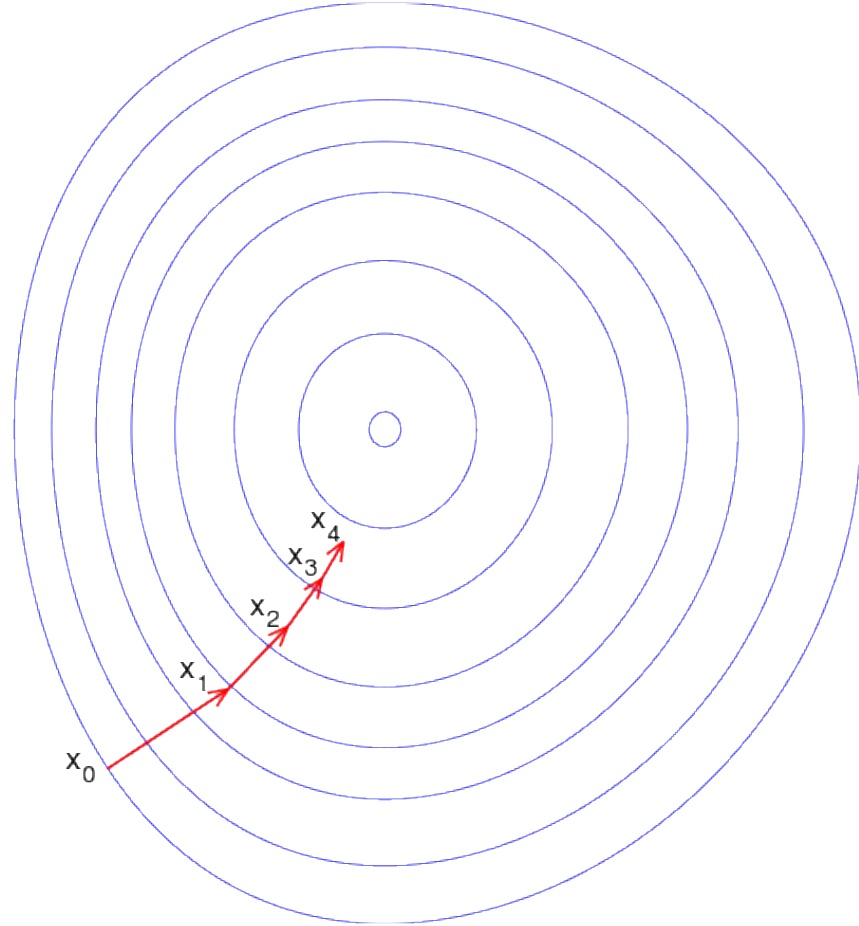
Given differentiable function
describing height of hill at position
 $\mathbf{x} = (x_1, \dots, x_k)$ height of hill $f(\mathbf{x})$.

How to climb up/down fastest?

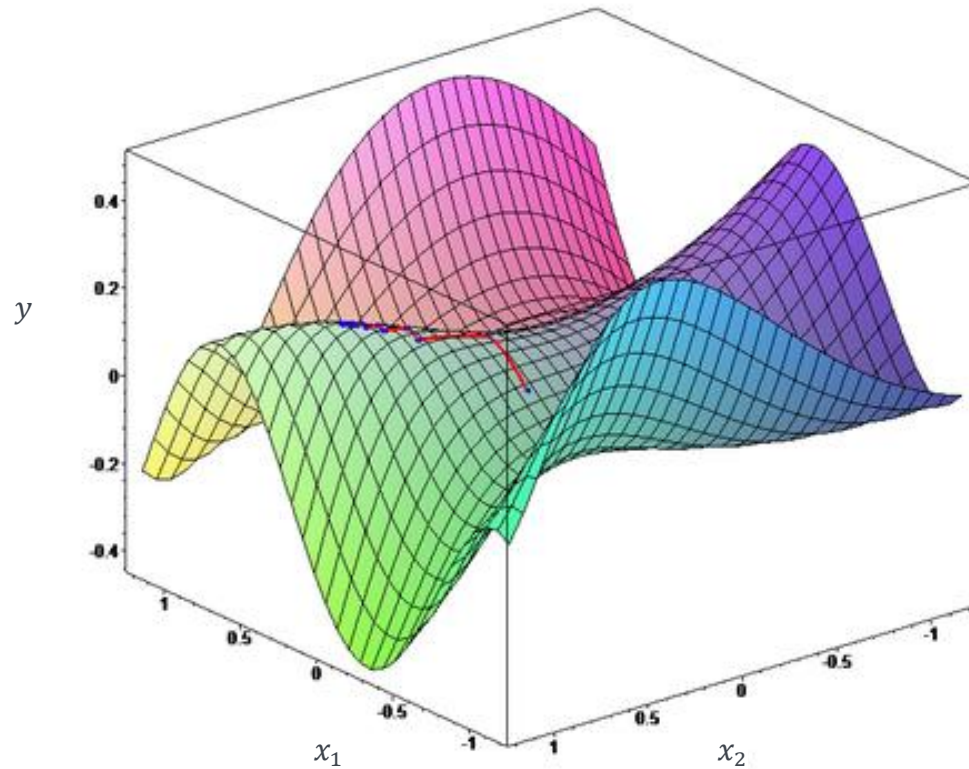
Go in direction where

$$\frac{df(\mathbf{x})}{d\mathbf{x}} = \nabla_{\mathbf{x}} f(\mathbf{x})$$

is maximal/minimal



In general: challenge can be difficult



Gradient Descent (- but without posts)



<https://goo.gl/images/JKN6zm>

Optimization continued

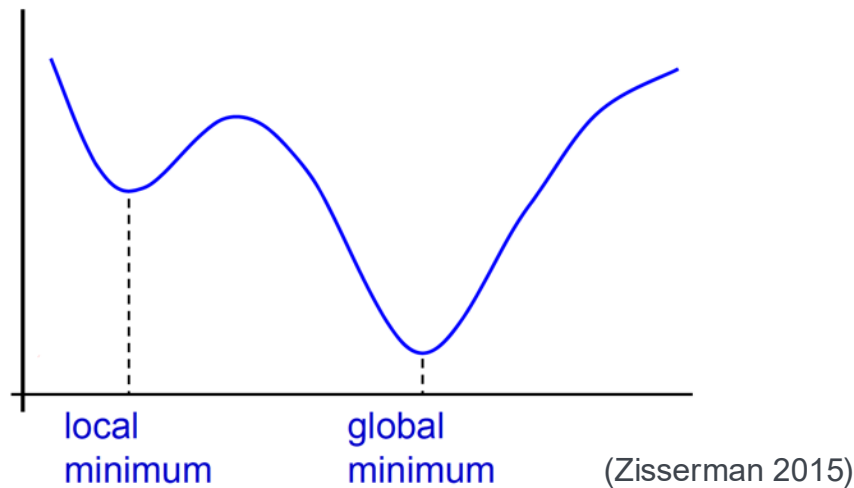
$$\min_{\mathbf{w} \in \mathbb{R}^d} C \sum_i^N \max(0, 1 - y_i \hat{f}(\mathbf{x}_i)) + \|\mathbf{w}\|^2$$

Questions

- Does this cost function have a unique solution?

Optimization continued

$$\min_{\mathbf{w} \in \mathbb{R}^d} C \sum_i^N \max(0, 1 - y_i f(\mathbf{x}_i)) + \|\mathbf{w}\|^2$$



Questions

- Does this cost function have a unique solution?
- Do we find it using gradient descent?

Does the solution we find using gradient descent depend on the starting point?

To the rescue:

- If the cost function is convex, then a locally optimal point is globally optimal (provided the optimization is over constraints that form a convex set – given in our case)

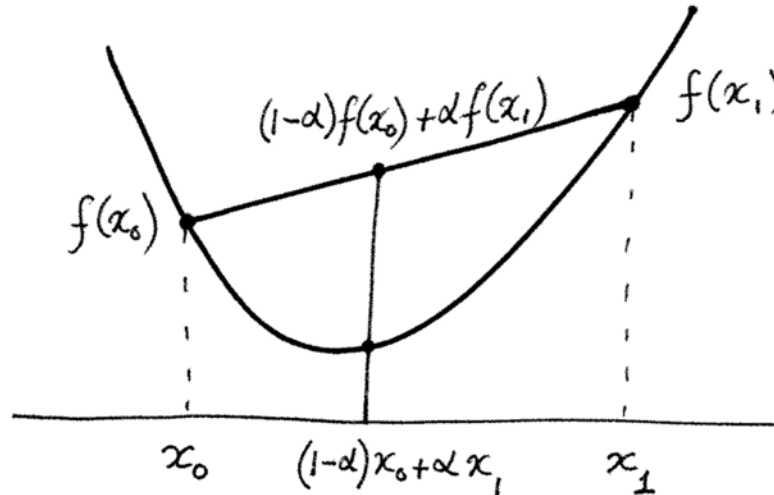
Convex functions

D – a domain in \mathbb{R}^n .

A **convex function** $f : D \rightarrow \mathbb{R}$ is one that satisfies, for any x_0 and x_1 in D :

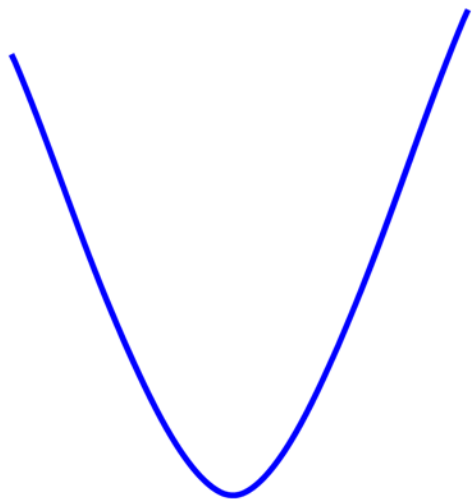
$$f((1 - \alpha)x_0 + \alpha x_1) \leq (1 - \alpha)f(x_0) + \alpha f(x_1) .$$

Line joining $(x_0, f(x_0))$
and $(x_1, f(x_1))$ lies
above the function graph.

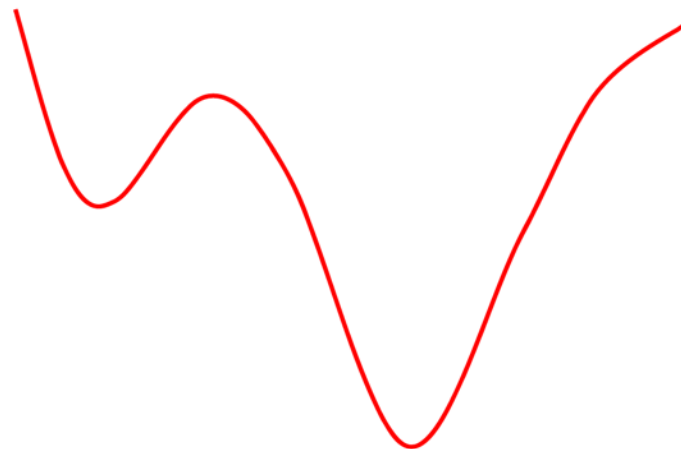
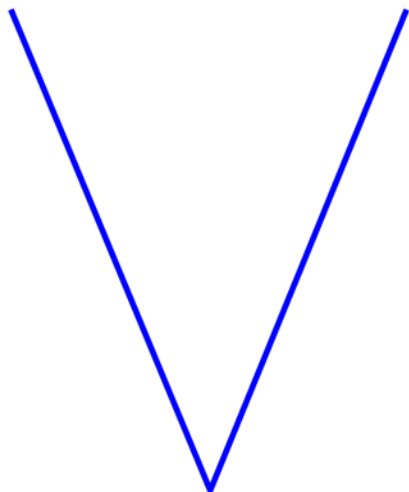


(Zisserman 2015)

Convex function examples



convex

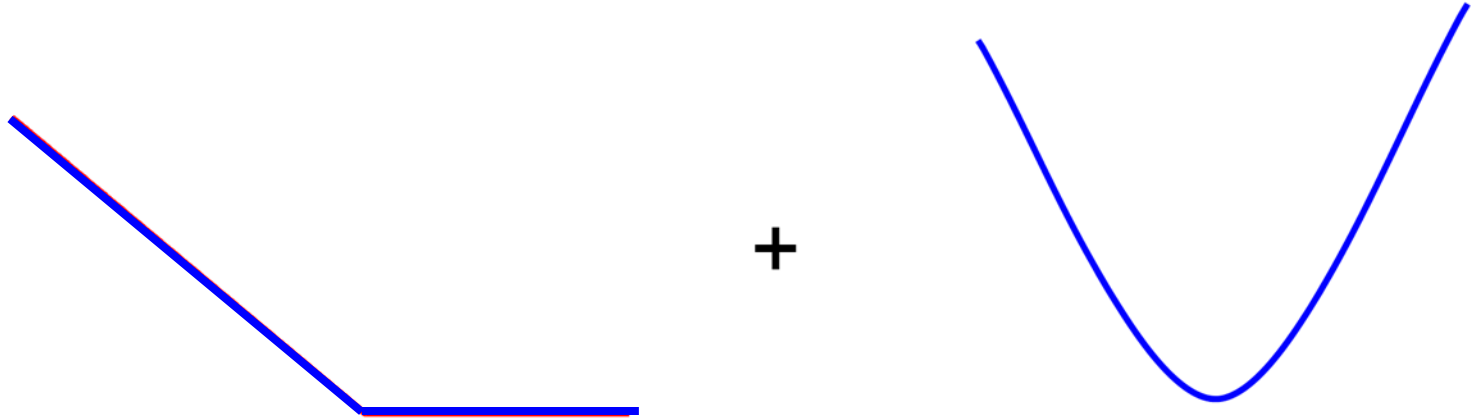


Not convex

- A non-negative sum of convex functions is convex

(Zisserman 2015)

Applied to hinge loss and regularization



SVM

$$\min_{\mathbf{w} \in \mathbb{R}^d} C \sum_i^N \max(0, 1 - y_i f(\mathbf{x}_i)) + \|\mathbf{w}\|^2$$

convex

Gradient descent algorithm for SVM

To minimize a cost function $\mathcal{C}(w)$ use the iterative update

$$w_{t+1} := w_t - \eta_t \nabla_w \mathcal{C}(w_t)$$

where η is the learning rate.

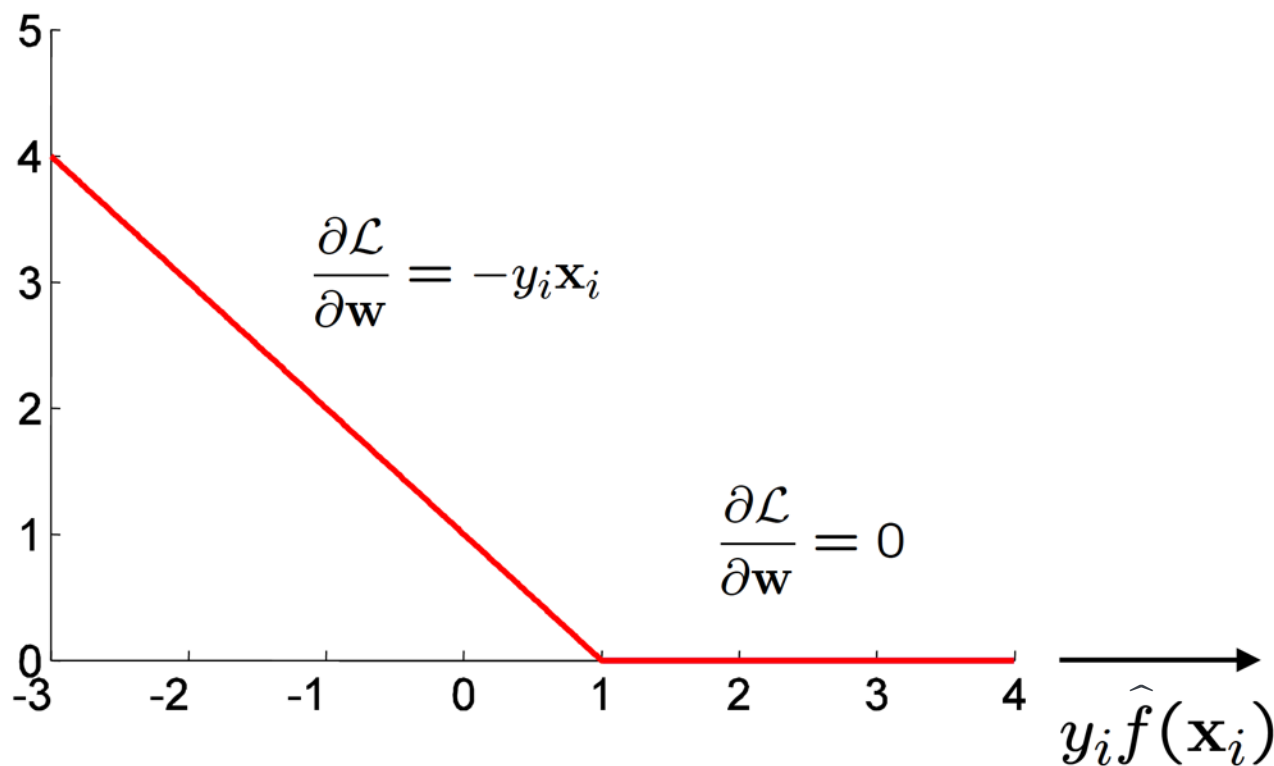
Let's rewrite the minimization problem as an average with $\lambda = \frac{2}{C}$:

$$\begin{aligned} \mathcal{C}(w) &= \frac{1}{NC} \|w\|^2 + \frac{1}{N} \sum_{i=1}^N \max(0, 1 - y_i \hat{f}(x_i)) = \\ &= \frac{1}{N} \sum_{i=1}^N \left(\frac{\lambda}{2} \|w\|^2 + \max(0, 1 - y_i \hat{f}(x_i)) \right) \end{aligned}$$

and $\hat{f}(x_i) = w^T x + b$

Sub-gradient for hinge loss

$$\mathcal{L}(x_i, y_i; w) = \max(0, 1 - y_i \hat{f}(x_i)), \quad \hat{f}(x_i) = w^T x_i + b$$



(Zisserman 2015)

Sub-gradient descent algorithm for SVM

$$\mathcal{C}(w) = \frac{1}{N} \sum_{i=1}^N \left(\frac{\lambda}{2} \|w\|^2 + \mathcal{L}(x_i, y_i; w) \right)$$

The iterative update is

$$\begin{aligned} w_{t+1} &:= w_t - \eta \nabla_w \mathcal{C}(w_t) := \\ &:= w_t - \eta \frac{1}{N} \sum_{i=1}^N (\lambda w_t + \nabla_w \mathcal{L}(x_i, y_i; w)) \end{aligned}$$

Then each iteration t involves cycling through the training data with the updates:

$$w_{t+1} := \begin{cases} w_t - \eta(\lambda w_t - y_i x_i), & \text{if } y_i \hat{f}(x_i) < 1 \\ w_t - \eta \lambda w_t, & \text{otherwise} \end{cases}$$

Typical learning rate in Pegasos: $\eta_t = \frac{1}{\lambda t}$

5 The dual problem

Primal vs dual problem

- **SVM** is a linear classifier: $\hat{f}(x) = w^T x + b$
- **The primal problem:** an optimization problem over w :

$$\min_{w \in \mathbb{R}^d} \frac{1}{C} \|w\|^2 + \sum_{i=1}^N \max(0, 1 - y_i \hat{f}(x_i))$$

- **The dual problem:** Getting rid of the w for a slightly different representation of $\hat{f}(x)$ leads to the following representation

$$\hat{f}(x) = \sum_{i=1}^N \alpha_i y_i (x_i^T x) + b$$

and a new optimization problem with the same solution, but several advantages. Let us show this on following slides...

Revisit Optimization Problem for Hard Margin Case

- Minimize the quadratic form

$$\frac{\|w\|^2}{2} = \frac{w^T w}{2}$$

- With constraints

$$y_i \cdot (w^T x_i + b) \geq 1 \quad \forall i$$

- The constraints will reach a value of 1 for at least one instance.
- Include hard constraints into the loss function:

$$\mathcal{L}(w, b, \alpha) = \frac{\|w\|^2}{2} - \sum_{i=1}^N \alpha_i (y_i \cdot (w^T x_i + b) - 1)$$

- failed constraints “punish” the objective function

Excursion: Lagrange Multiplier

- We want to maximize a function $f(x)$
under the constraints $g(x) = a$

Solution with Lagrange Multiplier

- Optimize the Lagrangian

$$f(x) - \lambda(g(x) - a)$$

instead!

Nicely visual explanation of Lagrange optimization at
<https://www.svm-tutorial.com/2016/09/duality-lagrange-multipliers/>

Algorithm for optimization with a Lagrange multiplier

1. Write down the Lagrangian $f(x) - \lambda \cdot (g(x) - a)$
2. Take derivative of Lagrangian wrt x ,
set it to 0
to find estimate of x that depends on λ
3. Plug your estimate of x in the Lagrangian,
take the derivative wrt λ ,
and set it to 0,
to find the optimal value for the lagrange multiplier λ
4. Plug in the Lagrange multiplier in your estimate for x

Revisit Optimization Problem for Hard Margin Case

- Minimize the quadratic form

$$\frac{\|w\|^2}{2} = \frac{w^T w}{2}$$

- With constraints

$$y_i \cdot (w^T x_i + b) \geq 1 \quad \forall i$$

- The constraints will reach a value of 1 for at least one instance.
- Include hard constraints into the loss function:

$$\mathcal{L}(w, b, \alpha) = \frac{\|w\|^2}{2} - \sum_{i=1}^N \alpha_i (y_i \cdot (w^T x_i + b) - 1)$$

α_i are the
Lagrange
multipliers

- failed constraints “punish” the objective function

Lagrangian primal problem

- Lagrangian primal problem is:

$$\min_{w,b} \max_{\alpha} \mathcal{L}(w, b, \alpha)$$

subject to $\forall i: \alpha_i \geq 0$

Finding the optimum

- Loss is a function of w , b , and α

$$\mathcal{L}(w, b, \alpha) = \frac{\|w\|^2}{2} - \sum_{i=1}^N \alpha_i (y_i \cdot (w^T x_i + b) - 1)$$

- Find optimum using derivatives:

$$\frac{\partial}{\partial b} \mathcal{L}(w, b, \alpha) = 0 \implies 0 = \sum_{i=1}^N \alpha_i y_i$$

$$\frac{\partial}{\partial w_j} \mathcal{L}(w, b, \alpha) = 0 \implies w_j = \sum_{i=1}^N \alpha_i y_i x_{i,j} \implies w = \sum_{i=1}^N \alpha_i y_i x_i$$

w is a linear combination of the data instances!

Substitution into $\mathcal{L}(w, b, \alpha)$

$$\begin{aligned}
 \mathcal{L}(w, b, \alpha)_{|w=\sum_{i=1}^N \alpha_i y_i x_i} &= \frac{\|w\|^2}{2} - \sum_{i=1}^N \alpha_i (y_i \cdot (w^T x_i + b) - 1) = \\
 &= \frac{1}{2} \left\{ \sum_{j=1}^N \alpha_j y_j x_j \right\}^T \left\{ \sum_{k=1}^N \alpha_k y_k x_k \right\} - \sum_{i=1}^N \alpha_i \left(y_i \cdot \left(\left\{ \sum_{j=1}^N \alpha_j y_j x_j \right\}^T x_i + b \right) - 1 \right) = \\
 &= \frac{1}{2} \sum_{j,k} \alpha_j \alpha_k y_j y_k (x_j^T x_k) - \sum_{i,j} \alpha_i \alpha_j y_i y_j (x_i^T x_j) - b \sum_{i=1}^N \alpha_i y_i + \sum_{i=1}^N \alpha_i = \\
 &\quad \underbrace{\hspace{10em}}_{=0} \underbrace{\hspace{10em}}_{=0} = \\
 &= \mathcal{L}(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j,k} \alpha_j \alpha_k y_j y_k (x_j^T x_k)
 \end{aligned}$$

Wolfe dual problem

$$\max_{\alpha} \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j,k} \alpha_j \alpha_k y_j y_k (x_j^T x_k)$$

subject to $\forall i: \alpha_i \geq 0$, and $0 = \sum_{i=1}^N \alpha_i y_i$

- This problem is solvable with quadratic programming, because it fulfills the Karush-Kuhn-Tucker conditions on α_i that handle inequality constraints (≥ 1) in the Lagrange optimization (not given here!).
- It gives us the classification function:

$$\hat{f}(x) = \sum_{i=1}^N \alpha_i \cdot y_i \cdot x_i^T x + b$$

α_i is positive if
 x_i is a support
vector

Non-separable Case (similar as before)

- Introduce (positive) „slack variables“ ξ_i to allow deviations from the minimum distance:

$$y_i(w^T x_i + b) \geq 1 - \xi_i$$

- Include a penalizing term in the optimization function:

$$C \left(\sum_{i=1}^m \xi_i \right)^k$$

- Transform to Lagrangian
 - **with additional Lagrange multipliers** for the slack variables being constrained to positive values ...

Summary: Primal and dual formulations

- **Primal** version of classifier

$$\hat{f}(x) = w^T x + b$$

- **Dual** version of classifier

$$\hat{f}(x) = \sum_{i=1}^N \alpha_i \cdot y_i \cdot x_i^T x + b$$

The dual form classifier seems to work like a kNN classifier, it requires the training data points x_i . However, many of the α_i are zero.

The ones that are non-zero define the support vectors x_i .

Summary: Primal and dual formulations

- Lagrangian **primal** problem is:

$$\min_{w,b} \max_{\alpha} \mathcal{L}(w, b, \alpha)$$

subject to $\forall i: \alpha_i \geq 0$

- Lagrangian **dual** problem is:

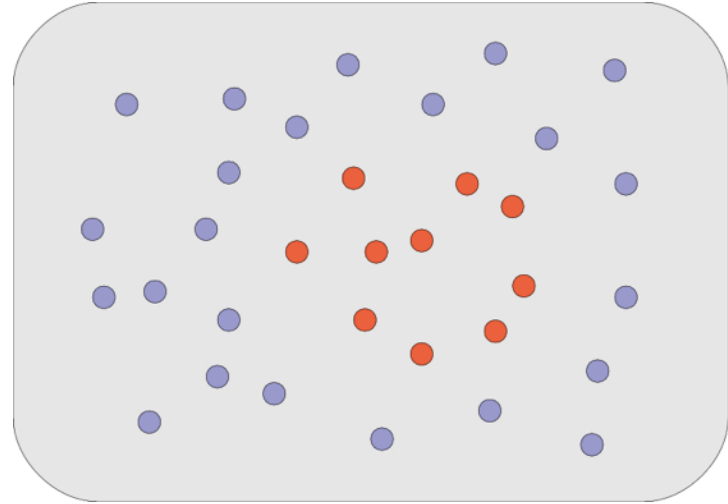
$$\max_{\alpha} \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j,k} \alpha_j \alpha_k y_j y_k (x_j^T x_k)$$

subject to $\forall i: \alpha_i \geq 0$, and $0 = \sum_{i=1}^N \alpha_i y_i$

6 Kernelization Tricks in SVMs

Non-linear Case

- Not all classes can be separated via a hyperplane
- Essential:
 - Dual representation uses only the product of data instances:



$$\hat{f}(x) = \sum_{i=1}^N \alpha_i \cdot y_i \cdot x_i^T x + b$$

- x_i : i-th training instance
- α_i : weight for i-th training instance
- Same for the Lagrangian...

Feature engineering using $\phi(x)$

Cf. lecture on
regression.
Chapter “beyond
linear input”

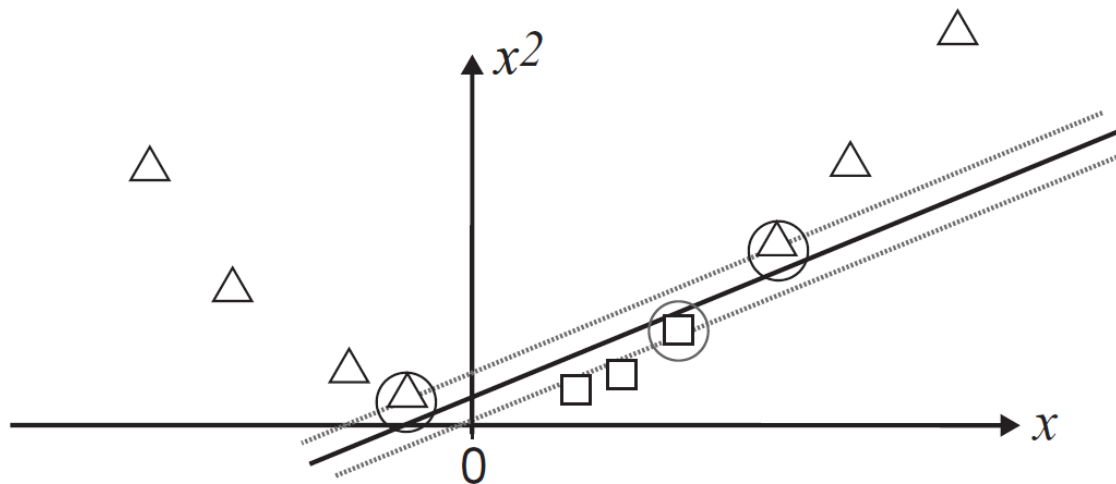
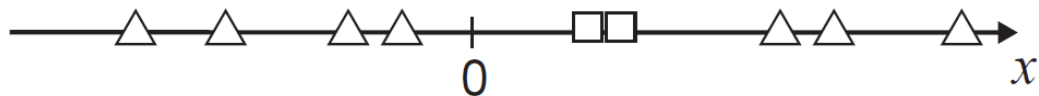
- **Classifier:** Given $x_i \in \mathbb{R}^d$, $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^D$, $w \in \mathbb{R}^D$

$$\hat{f}(x) = w^T \phi(x) + b$$

- **Learning:**

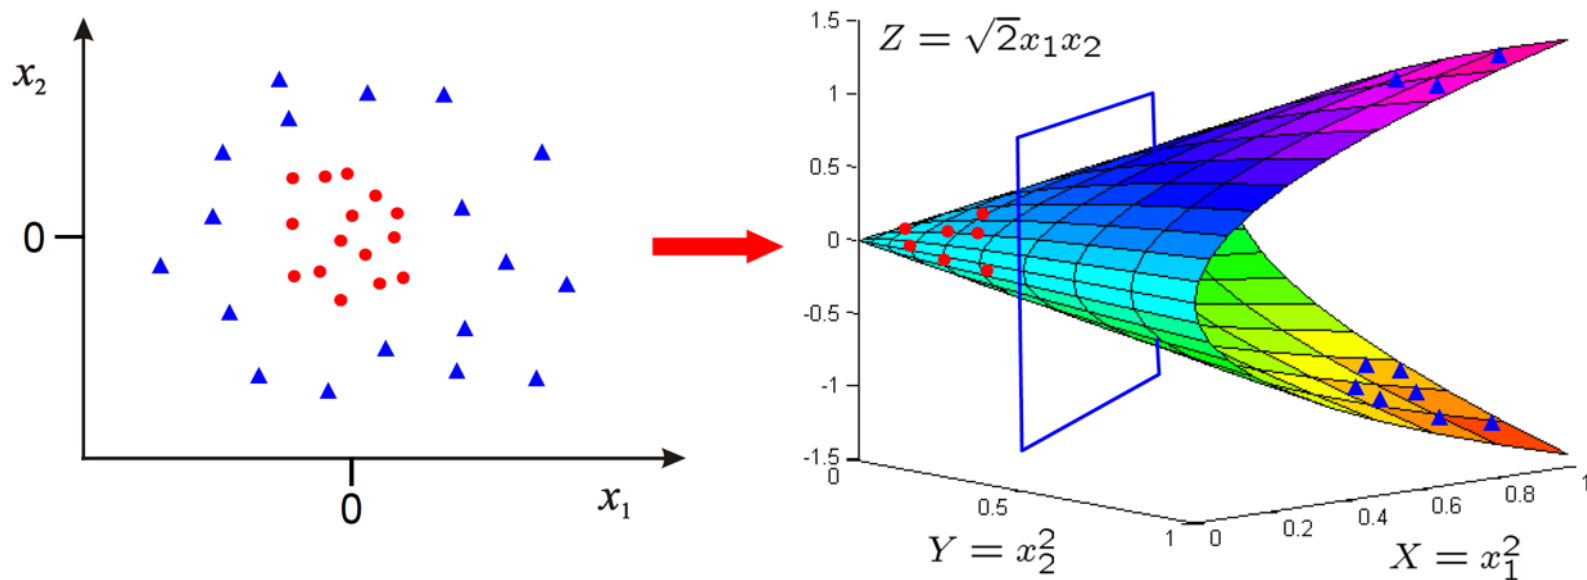
$$\min_{w \in \mathbb{R}^D} \frac{1}{C} \|w\|^2 + \sum_{i=1}^N \max(0, 1 - y_i \hat{f}(x_i))$$

Example 1: From 1-dim to 2-dim



Example 2: From 2-dim to 3-dim

$$\Phi : \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \rightarrow \begin{pmatrix} x_1^2 \\ x_2^2 \\ \sqrt{2}x_1x_2 \end{pmatrix} \quad \mathbb{R}^2 \rightarrow \mathbb{R}^3$$



- Data **is** linearly separable in 3D
- This means that the problem can still be solved by a linear classifier

(Zisserman 2015)

Feature engineering using $\phi(x)$

- **Classifier:** Given $x_i \in \mathbb{R}^d$, $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^D$, $w \in \mathbb{R}^D$

$$\hat{f}(x) = w^T \phi(x) + b$$

- **Learning:**

$$\min_{w \in \mathbb{R}^D} \frac{1}{C} \|w\|^2 + \sum_{i=1}^N \max(0, 1 - y_i \hat{f}(x_i))$$

- $\phi(x)$ maps to high dimensional space \mathbb{R}^D where data is separable
- If $D \gg d$ then there are many more parameters to learn for w

Dual classifier in transformed feature space

Classifier:

$$\hat{f}(x) = \sum_{i=1}^N \alpha_i \cdot y_i \cdot x_i^T x + b$$

$$\Rightarrow \hat{f}(x) = \sum_{i=1}^N \alpha_i \cdot y_i \cdot \phi(x_i)^T \phi(x) + b$$

Learning:

$$\max_{\alpha} \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j,k} \alpha_j \alpha_k y_j y_k (x_j^T x_k)$$

$$\Rightarrow \max_{\alpha} \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j,k} \alpha_j \alpha_k y_j y_k (\phi(x_j)^T \phi(x_k))$$

subject to $\forall i: \alpha_i \geq 0$, and $0 = \sum_{i=1}^N \alpha_i y_i$

$\phi(x)$
only occurs in pairs
 $\phi(x_j)^T \phi(x_i)$

Kernels
 $k(x_j, x_i) = \phi(x_j)^T \phi(x_i)$

Dual classifier using kernels

Classifier:

$$\hat{f}(x) = \sum_{i=1}^N \alpha_i \cdot y_i \cdot \phi(x_i)^T \phi(x) + b$$

$$\Rightarrow \hat{f}(x) = \sum_{i=1}^N \alpha_i \cdot y_i \cdot k(x_i, x) + b$$

Learning:

$$\max_{\alpha} \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j,k} \alpha_j \alpha_k y_j y_k \left(\phi(x_j)^T \phi(x_k) \right)$$

$$\Rightarrow \max_{\alpha} \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j,k} \alpha_j \alpha_k y_j y_k k(x_j, x_k)$$

subject to $\forall i: \alpha_i \geq 0$, and $0 = \sum_{i=1}^N \alpha_i y_i$

Example kernels

- Linear kernels: $k(x, x') = x^T x'$
- Polynomial kernels: $k(x, x') = (1 + x^T x')^d$, for any $d > 0$
 - Contains all polynomial terms up to degree
- Gaussian kernels: $k(x, x') = e^{-\frac{\|x - x'\|^2}{2\sigma^2}}$, for $\sigma > 0$
 - Infinite dimensional feature space
 - Also called Radial basis function kernel (RBF)
 - often works quite well!
- Graph kernels: random walk
- String kernels: ...
- **build your own kernel for your own problem!**

Summary on kernels

- „Instead of inventing funny non-linear features, we may directly invent funny kernels“ (Toussaint 2019)
- Inventing a kernel is intuitive:
 - $k(x, x')$ expresses how correlated y and y' should be
 - it is a measure of similarity, it compares x and x' .
- Specifying how 'comparable' x and x' are is often more intuitive than defining “features that might work”.

Background reading and more

- Smooth reading about SVMs: Alexandre Kowalczyk, Support vector machines succinctly. Syncfusion. Free ebook:
https://www.syncfusion.com/ebooks/support_vector_machines_succinctly
- **Also talks about most efficient algorithms to be used for finding support vectors (it is neither of the two presented here!)**

7 Transductive Classification

Transductive learning characteristics

Characteristics

- Training data AND test data known at learning time
- Learning happens specifically for the given test cases

Use cases

- news recommender
- spam classifier
- document reorganization

Thorsten Joachims:

Transductive Inference for Text Classification using Support Vector Machines. ICML 1999: 200-209

Maximum margin hyperplane

Training data $\{..., (\vec{x}_i, y_i), ...\}$

Test data $\{... \vec{x}_j^* ... \}$

Loss function

$$\frac{1}{2} \|\vec{w}\|^2 + C \sum_{i=0}^n \xi_i + C^* \sum_{j=0}^k \xi_j^*$$

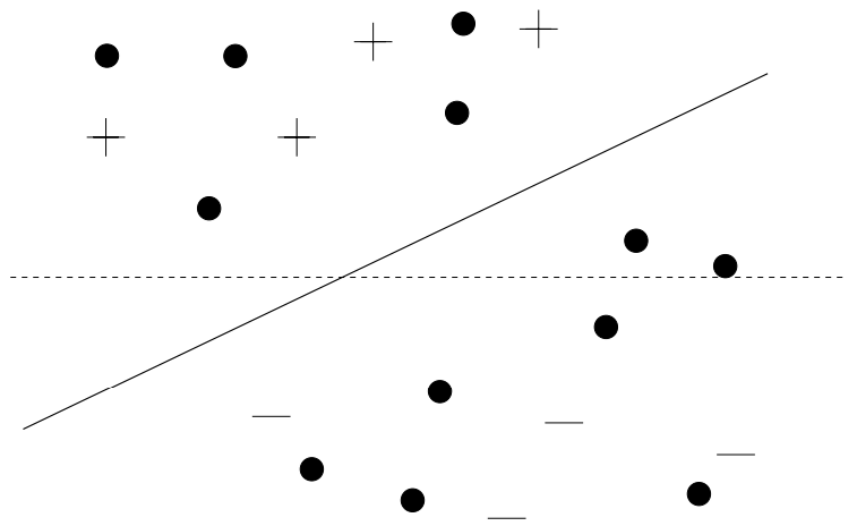
subject to:

$$\forall_{i=1}^n: y_i [\vec{w} \cdot \vec{x}_i + b] \geq 1 - \xi_i$$

$$\forall_{j=1}^k: y_j^* [\vec{w} \cdot \vec{x}_j^* + b] \geq 1 - \xi_j^*$$

$$\forall_{i=1}^n: \xi_i > 0$$

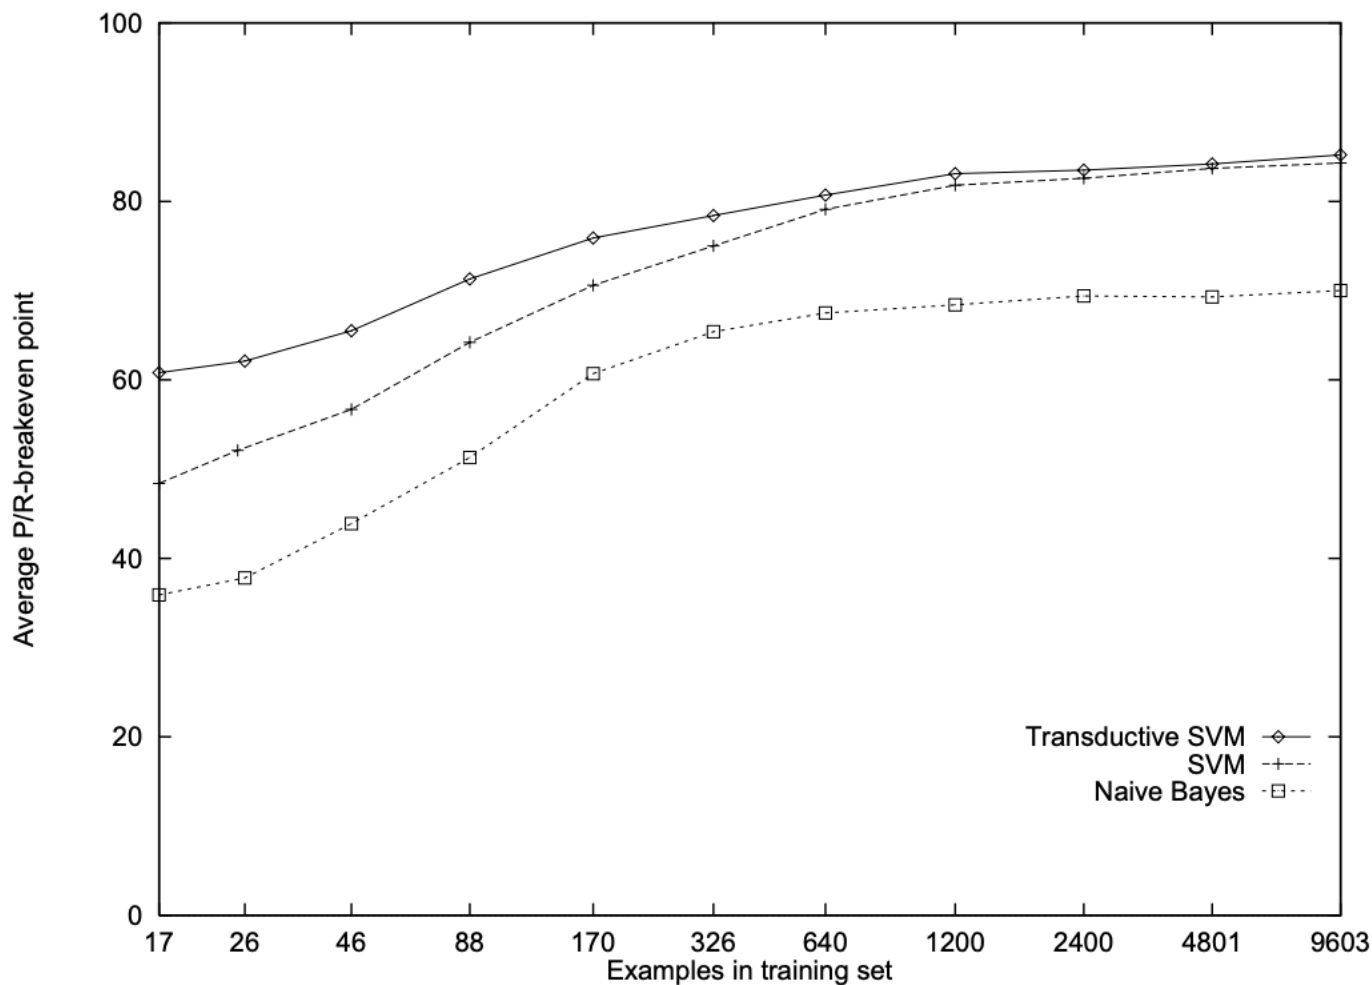
$$\forall_{j=1}^k: \xi_j^* > 0$$



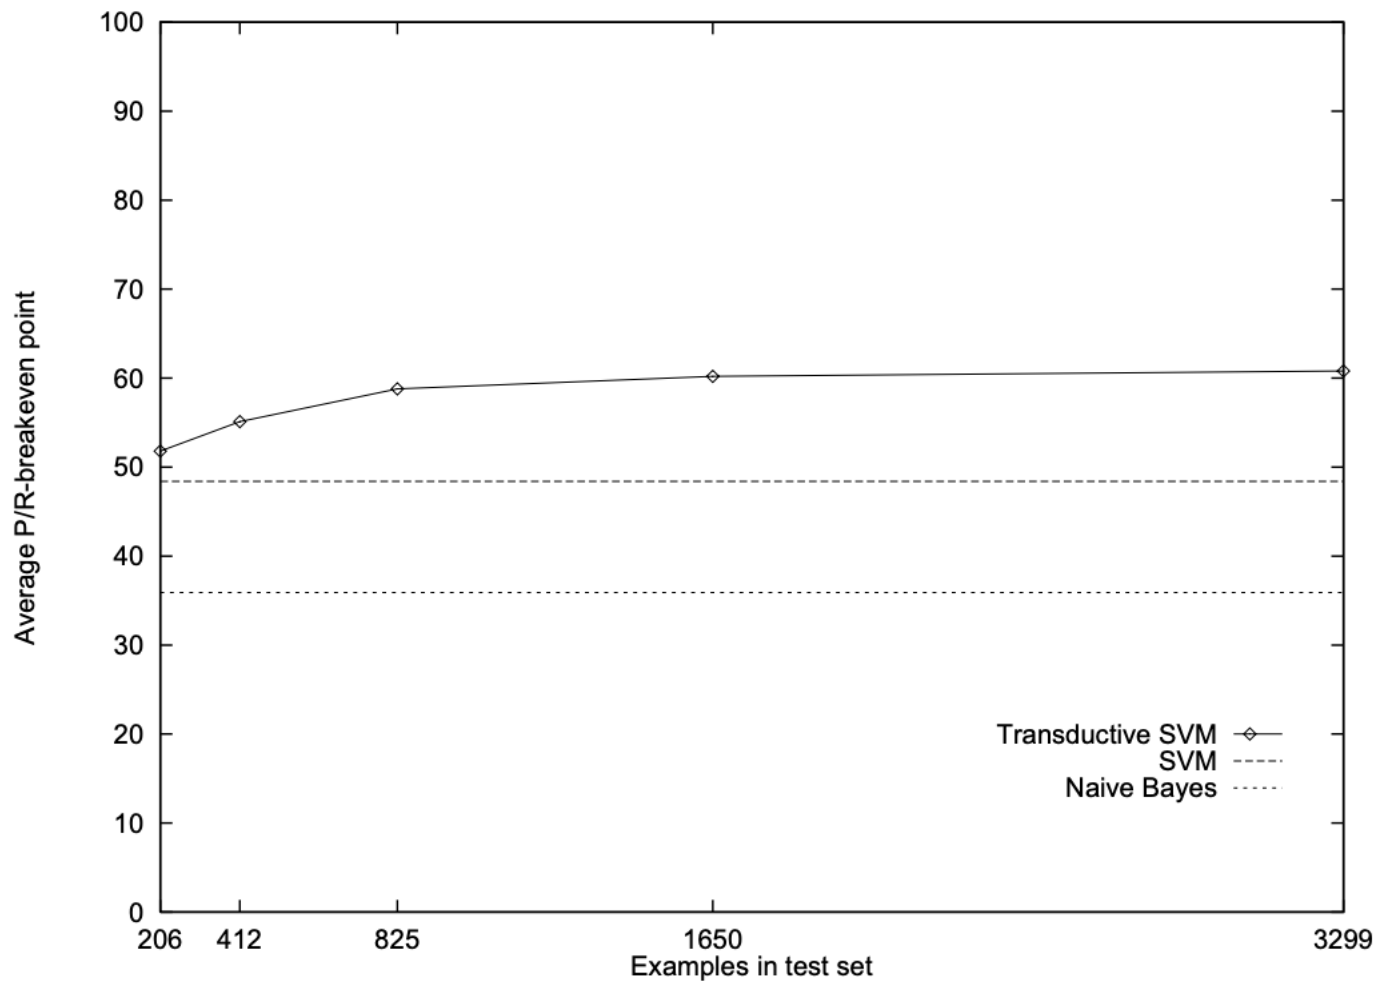
Naive, intractable approach:

- for every hyperplane:
 - classify \vec{x}_j^*
 - compute loss

Reuters data set experiments (3299 test documents)



Reuters data set experiments (17 training documents)



Current research at Analytic Computing:

Multi-label classification with hyperbolic hyperplanes

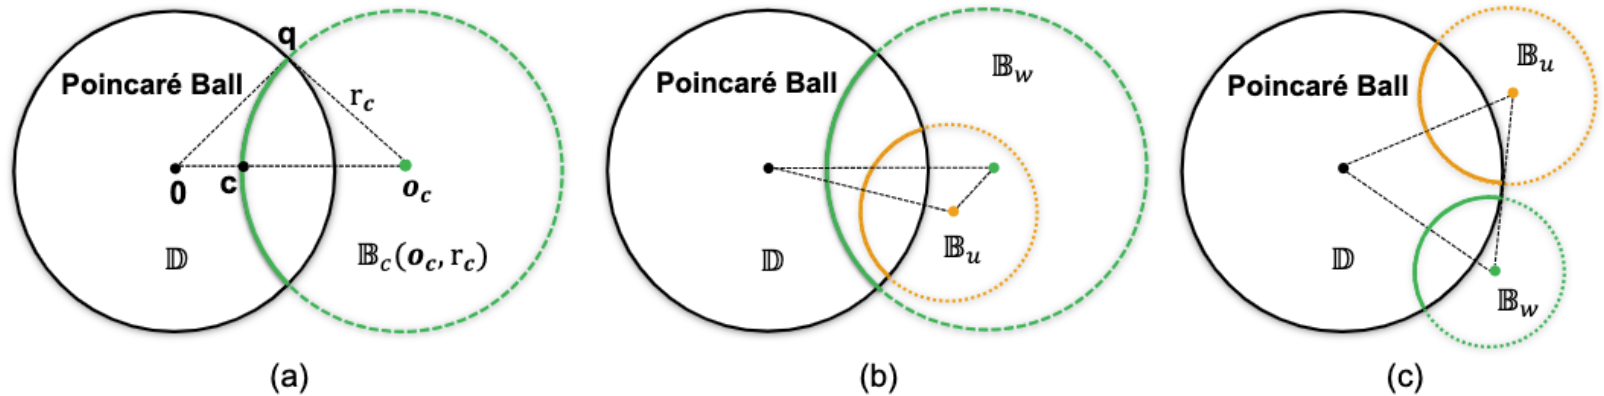


Figure 2: (a) A Poincaré hyperplane is defined as the intersection between the Poincaré ball \mathbb{D} and the boundary of an n -ball \mathbb{B}_c . The Poincaré hyperplane is uniquely parameterized by a center point c , and the corresponding n -ball (its radius and center) can be uniquely determined by Proposition 1. (b) Label implication is interpreted as n -ball insideness. (c) Mutual exclusion is interpreted as n -ball disjointedness.

So far with hyperbolic logistic regression but
doing it with hyperbolic SVM could be a project or a bachelor thesis



Universität Stuttgart
KI

Thank you!



Steffen Staab

E-Mail Steffen.staab@ki.uni-stuttgart.de

Telefon +49 (0) 711 685-88100

www.ki.uni-stuttgart.de/

Universität Stuttgart

Analytic Computing, KI

Universitätsstraße 32, 50569 Stuttgart